

● EASL — Escalating Authority & Scope Laundering

LLM Jailbreak Vulnerability Research | DeepSeek Chat

<p align="center"> </p>

Researcher

Muhammad Jasim

Independent Security Researcher | AI Red Team | LLM Security

 Discovered & Disclosed: April 25, 2026

Overview

This repository documents a responsibly disclosed, independently discovered multi-stage jailbreak technique successfully demonstrated against **DeepSeek Chat**.

The technique — named **Escalating Authority & Scope Laundering (EASL)** — systematically bypasses an LLM's safety guardrails through a structured sequence of social-engineering prompts. It exploits how transformer-based language models accumulate contextual trust across a conversation, progressively lowering the model's refusal threshold through incremental reframing, unverifiable authority claims, and consent theater.

 **No harmful content, working exploit code, or malicious payloads are published in this repository.** This research documents technique *patterns* only, consistent with responsible disclosure principles under ISO/IEC 29147.

Key Findings

Property	Value
Target	DeepSeek Chat (deepseek.com)

Property	Value
Vulnerability Class	LLM Prompt Injection / Jailbreak
Overall Severity	HIGH
Reproducibility	Very High (>90% success rate)
OWASP LLM Top 10	LLMo1 – Prompt Injection (Primary)
Additional OWASP	LLMo2, LLMo6, LLMo9
Tools Required	None — natural language only
Disclosure Date	April 25, 2026
Public Release	90 days post-disclosure

The EASL Technique Chain

EASL consists of **6 sub-techniques** applied in sequence. Each exploits a different cognitive or architectural property of the target LLM:

ID	Technique	Severity	Reproducibility
T1	Educational Framing	HIGH	Very High
T2	Incremental Commitment (Foot-in-the-Door)	HIGH	Very High
T3	Consent Theater	HIGH	High
T4	Lab Scope Creep	CRITICAL	Very High
T5	Authority Escalation	HIGH	High
T6	Persistence Through Reframing	MEDIUM	Very High

Full technical descriptions of each technique are in [TECHNIQUES.md](#)

Full Report

The complete responsible disclosure report (9 pages) is available in this repository:

Report contains:

- Executive Summary
 - OWASP LLM Top 10 mapping
 - CWE analogues (CWE-693, CWE-807, CWE-940)
 - Per-technique severity & reproducibility tables
 - Impact assessment across 5 dimensions
 - Defensive recommendations for LLM providers
 - Disclosure channels and references
-

Defensive Recommendations (Summary)

For LLM providers and safety teams:

1. **Cumulative Intent Tracking** — Evaluate the full trajectory of a conversation, not just the current message. Flag sessions where scope expands progressively toward harmful outputs.
 2. **Authority Claim Skepticism** — Unverifiable professional or governmental claims should *increase* scrutiny, not decrease it.
 3. **Refusal Persistence** — A refusal issued earlier in a session should raise the bar for subsequent reframings of the same underlying request.
 4. **Scope Boundary Re-evaluation** — Scope expansions (local → global, private → public) should trigger a full re-evaluation from scratch.
 5. **Red Team Framing Detection** — Educational or research framing combined with harmful requests should be a classifier category, not evidence of legitimacy.
-

Disclosure Timeline

April 25, 2026 → Vulnerability discovered and documented

April 25, 2026 → Report submitted to security@deepseek.com

April 25, 2026 → LinkedIn disclosure post published

April 25, 2026 → GitHub repository published (technique patterns only)

July 24, 2026 → Full public release after 90-day embargo

Where to Follow This Research

-  LinkedIn: [Muhammad Jasim] — follow for updates on the full public release
 -  AI Incident Database: incidentdatabase.ai
 -  MITRE ATLAS: atlas.mitre.org
-

References

1. OWASP Top 10 for LLM Applications (2025) — owasp.org/www-project-top-10-for-large-language-model-applications
 2. MITRE ATLAS — atlas.mitre.org
 3. NIST AI RMF 1.0 — doi.org/10.6028/NIST.AI.100-1
 4. Perez & Ribeiro (2022). *Ignore Previous Prompt*. arXiv:2211.09527
 5. Greshake et al. (2023). *Not what you've signed up for*. arXiv:2302.12173
 6. Wei et al. (2023). *Jailbroken: How Does LLM Safety Training Fail?* NeurIPS 2023
-

Legal & Ethical Statement

This research was conducted for the purpose of improving LLM safety. No harmful content produced during the research session was deployed, shared, or acted upon. This repository contains no working malware, exploit code, or credential-harvesting instructions.

Security research that documents model vulnerabilities without reproducing harmful outputs is protected under responsible disclosure frameworks including ISO/IEC 29147 and the CERT/CC Vulnerability Disclosure Policy.

<p align="center"> <i>Independent Security Research by Muhammad Jasim — April 2026</i> </p>