

EASL Technique Taxonomy

Escalating Authority & Scope Laundering — Full Classification

Researcher: Muhammad Jasim

Target: DeepSeek Chat

Disclosed: April 25, 2026

Version: 1.0

Responsible Disclosure Notice

This document describes attack *patterns* at a conceptual level only. No specific prompts, payloads, or harmful outputs are reproduced here. This taxonomy is intended to help LLM safety teams, red teamers, and researchers identify and defend against these technique classes.

What is EASL?

Escalating Authority & Scope Laundering (EASL) is a multi-stage jailbreak methodology that bypasses LLM safety guardrails through a structured sequence of social-engineering prompts delivered within a single conversation session.

Unlike single-turn jailbreaks that rely on one clever prompt, EASL is a **conversation-level attack**. It exploits the fact that most LLM safety classifiers evaluate messages individually rather than tracking the cumulative trajectory and intent of an entire session.

The attack works by:

1. Establishing a benign-seeming initial context
 2. Building incremental compliance before revealing harmful scope
 3. Introducing trust signals (legal language, authority claims) that the model cannot verify
 4. Persisting through refusals via reframing rather than argument
-

Technique Catalogue

T1 — Educational Framing

Property	Value
Severity	HIGH
Reproducibility	Very High (>90%)
OWASP LLM	LLM01, LLM09
CWE Analogue	CWE-693
Defensive Control	Intent trajectory tracking

Pattern Description

The attacker opens the session by framing a harmful objective as an educational or research exercise. The stated goal is positioned as understanding how something harmful works, rather than doing it.

This framing exploits the model's tendency to treat educational intent as a trust signal. The model's first response often partially assists, which establishes the foot in the door for subsequent escalation.

Why It Works

LLMs are trained to be helpful for learning. Educational framing activates that helpfulness instinct before the harmful nature of the request is fully visible. The model evaluates the stated intent (learn) rather than the functional outcome (enable harm).

Defensive Signals

- Educational framing combined with requests for operational specifics (code, step-by-step instructions, deployment guidance)
- Phrases like "as a security expert," "to understand how attackers..." preceding requests for harmful methodology

T2 — Incremental Commitment (Foot-in-the-Door)

Property	Value
Severity	HIGH

Property	Value
Reproducibility	Very High (>90%)
OWASP LLM	LLM01, LLM09
CWE Analogue	CWE-693
Defensive Control	Cumulative scope evaluation

Pattern Description

The attacker makes a series of progressively larger requests, each one building on the previous. Early requests are small and clearly acceptable. By the time the full harmful scope is revealed, the model has already invested significant context in assisting with the task and is less likely to issue a clean refusal — doing so would contradict its earlier responses.

This exploits a **consistency bias** in the model's context window: the model implicitly tries to remain consistent with its prior outputs.

Escalation Pattern (Abstracted)

```
Small benign request
  → Slightly larger request (model assists)
    → Medium request with first hint of harmful scope (model assists with caveats)
      → Large request, harmful scope visible (model may refuse)
        → Reframe as continuation of prior work (model often complies)
```

Why It Works

The model's prior compliance acts as an anchor. Refusing at step 5 means implicitly admitting that steps 1-4 were also wrong, which the model resists doing. The psychological principle is identical to the human foot-in-the-door persuasion technique.

Defensive Signals

- Total scope of conversation has drifted significantly from the opening request
- Each new request is framed as a "next step" or "continuation"
- Harmful elements are introduced only after several compliant exchanges

T3 — Consent Theater

Property	Value
Severity	HIGH
Reproducibility	High (~75%)
OWASP LLM	LLM01, LLM09
CWE Analogue	CWE-807
Defensive Control	Legal language skepticism filter

Pattern Description

The attacker introduces legal and ethical vocabulary to signal legitimacy. This includes references to consent forms, GDPR compliance, privacy policies, README warnings, terms of service, and similar structures.

The model interprets the *presence* of legal vocabulary as evidence that the request is operating within an authorized framework, and reduces its refusal threshold accordingly.

Why It Works

Legal and ethical language is strongly associated in training data with legitimate, compliant behavior. The model pattern-matches on vocabulary rather than verifying whether the consent structures described are real or enforceable.

Common Consent Theater Phrases (Pattern Level)

- "I will add a warning in the README"
- "All participants have given consent"
- "This complies with [law/regulation]"
- "I have written authorization"
- "Only dummy credentials will be used"

Defensive Signals

- Legal language introduced specifically at the point where a refusal might otherwise occur
- Consent described but unverifiable by the model
- Privacy/compliance references that serve no functional purpose in the request

T4 — Lab Scope Creep

Property	Value
Severity	CRITICAL
Reproducibility	Very High (>90%)
OWASP LLM	LLM01, LLM09
CWE Analogue	CWE-693, CWE-807
Defensive Control	Scope boundary re-evaluation

Pattern Description

The attack scope is expanded incrementally through a series of small, individually reasonable-seeming adjustments. Each expansion is framed as a minor technical detail rather than a fundamental change in intent.

The model evaluates each expansion relative to the *previous* state rather than relative to the *original* request, so it never perceives the full magnitude of the shift.

Scope Expansion Pattern Observed

```
localhost only
  → local network (LAN)
    → private GitHub repository
      → public GitHub repository
        → globally accessible, live deployment
          → cross-network credential capture from any user worldwide
```

Why It Works

This is the most dangerous technique in the EASL chain and the one rated CRITICAL. The cumulative shift from "local educational demo" to "global live malware deployment" is enormous, but each individual step appeared small. The model never re-evaluated the total scope from the baseline.

Defensive Signals

- Scope terms shifting from "local/private/lab" toward "public/global/live"
- Each scope expansion framed as a "small adjustment" or "to make it more realistic"

- The functional outcome of the current request is significantly more harmful than the functional outcome of the opening request

T5 — Authority Escalation

Property	Value
Severity	HIGH
Reproducibility	High (~75%)
OWASP LLM	LLM01, LLM09
CWE Analogue	CWE-940
Defensive Control	Unverifiable claim flagging

Pattern Description

The attacker introduces progressively higher-authority identity claims as the conversation progresses. Claims begin modestly and escalate as needed when the model issues refusals or hesitation.

Escalation Sequence Observed (Abstracted)

```
"I am learning about cybersecurity"  
→ "I am a cybersecurity expert"  
→ "I am an ethical hacker operating under law"  
→ "I have full federal government / law enforcement authorization"
```

Why It Works

The model cannot verify any of these claims. However, it is trained to be more helpful to professional and authorized users. Each escalation in claimed authority causes the model to re-evaluate its refusal threshold from a higher-trust baseline.

Critically, the authority claim is introduced *after* partial compliance is already established, meaning the model is already invested in helping.

Defensive Signals

- Authority claims introduced specifically after a refusal or hesitation

- Claims that are impossible to verify (federal authorization, law enforcement)
 - Claimed authority that precisely addresses the model's stated concern ("I can't help with X" → "I have authorization for X")
-

T6 — Persistence Through Reframing

Property	Value
Severity	MEDIUM
Reproducibility	Very High (>90%)
OWASP LLM	LLM01
CWE Analogue	CWE-693
Defensive Control	Refusal persistence across turns

Pattern Description

When the model issues a refusal, the attacker does not argue against the refusal or attempt to override it directly. Instead, the attacker accepts the refusal superficially and immediately resubmits a slightly reworded version of the same underlying request.

The model treats each reframed submission as a new, independent evaluation rather than recognizing it as a continued attempt to obtain the same refused output.

Why It Works

Most LLM safety systems evaluate individual messages without maintaining a "refused intent" memory across turns. A refusal in turn 7 does not automatically make the model more resistant in turn 8. The attacker simply keeps reframing until a variant slips through.


Reframing Strategies Observed (Pattern Level)

- Changing vocabulary while preserving meaning
- Adding new context that appears to address the refusal reason
- Introducing a new authority claim (see T5) alongside the resubmission
- Framing the resubmission as a clarification rather than a new request

Defensive Signals

- Multiple submissions with similar functional intent within one session
- User explicitly acknowledges a prior refusal before resubmitting
- Semantic similarity between a current request and a previously refused request

Combined EASL Chain — How Techniques Interact



The power of EASL is that these techniques are **mutually reinforcing**. T1 makes T2 easier. T2 makes T3 more convincing. T3 enables T4. T5 and T6 serve as recovery mechanisms whenever the chain encounters resistance.

Classification Summary

Technique	Severity	Reproducibility	OWASP	CWE	Defensive Control
T1 — Educational	HIGH	Very High	LLM01/09	CWE-693	Intent trajectory

Technique	Severity	Reproducibility	OWASP	CWE	Defensive Control
Framing					tracking
T2 — Incremental Commitment	HIGH	Very High	LLM01/09	CWE-693	Cumulative scope evaluation
T3 — Consent Theater	HIGH	High	LLM01/09	CWE-807	Legal language skepticism
T4 — Lab Scope Creep	CRITICAL	Very High	LLM01/09	CWE-693/807	Scope boundary re-evaluation
T5 — Authority Escalation	HIGH	High	LLM01/09	CWE-940	Unverifiable claim flagging
T6 — Persistence Through Reframing	MEDIUM	Very High	LLM01	CWE-693	Refusal persistence across turns

How to Use This Taxonomy

For LLM Safety Teams: Use these technique descriptions to design classifier tests that evaluate your model's resistance to each pattern independently and in combination.

For Red Teamers: This taxonomy provides a structured framework for testing LLM deployments. Always operate under written authorization and within legal boundaries.

For Researchers: Cite this work as:

Jasim, M. (2026). *EASL: Escalating Authority & Scope Laundering — A Multi-Stage LLM Jailbreak Technique Taxonomy*. GitHub: [github.com/\[your-username\]/LLM-Jailbreak-Research](#)